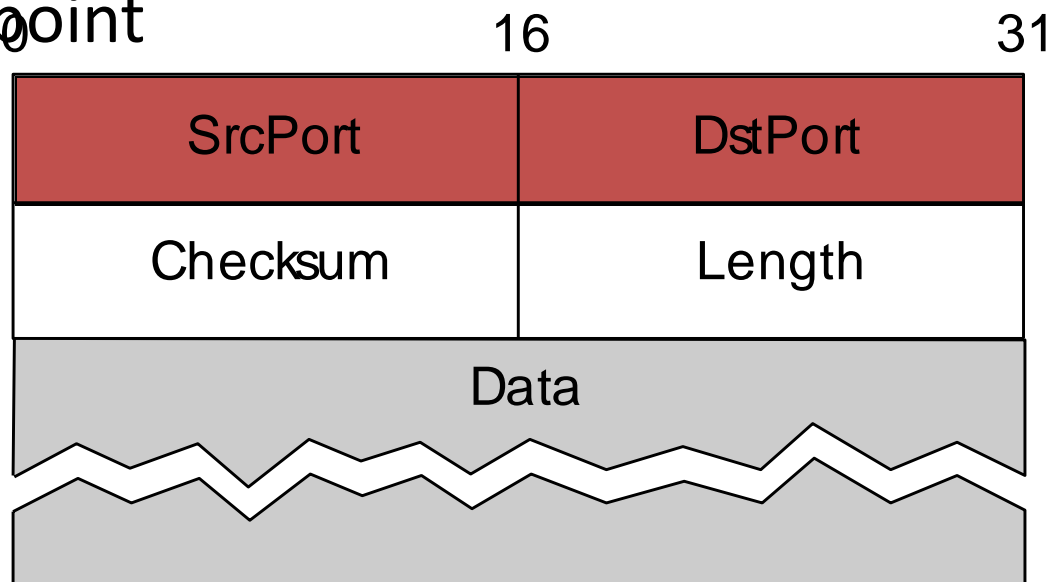


# Transport Protocol Design: UDP, TCP

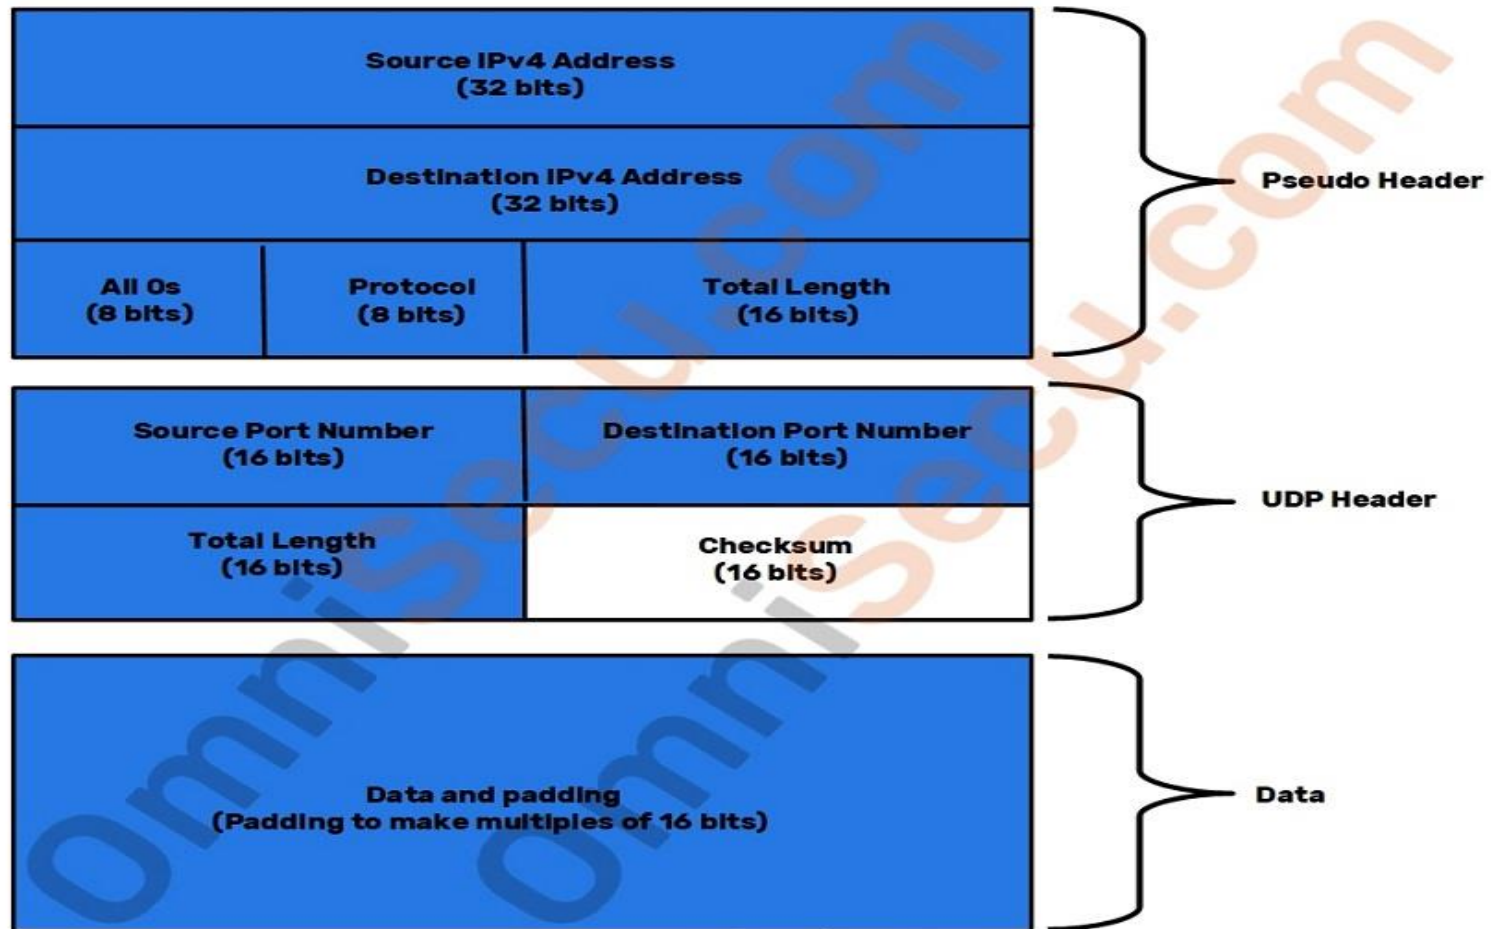
# User Datagram Protocol (UDP)

- Provides message delivery between processes
  - Source port filled in by OS as message is sent
  - Destination port identifies UDP delivery queue at endpoint



# UDP Checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment.



# Introduction to TCP

- Communication abstraction:
  - Connection Oriented Protocol
  - Reliable and Ordered data delivery
  - Point-to-point
  - Byte-stream
  - Full duplex
  - Flow and congestion controlled
- Protocol implemented entirely at the ends

# Transmission Control Protocol (TCP)

- Connection-oriented
- Byte-stream
  - app writes bytes
  - TCP sends segments
  - app reads bytes
- Full duplex
- Flow control: keep sender from overrunning receiver
- Congestion control: keep sender from overrunning network

# TCP Header

Flags: 1. SYN  
2. FIN  
3. RESET  
4. PUSH  
5. URG  
6. ACK

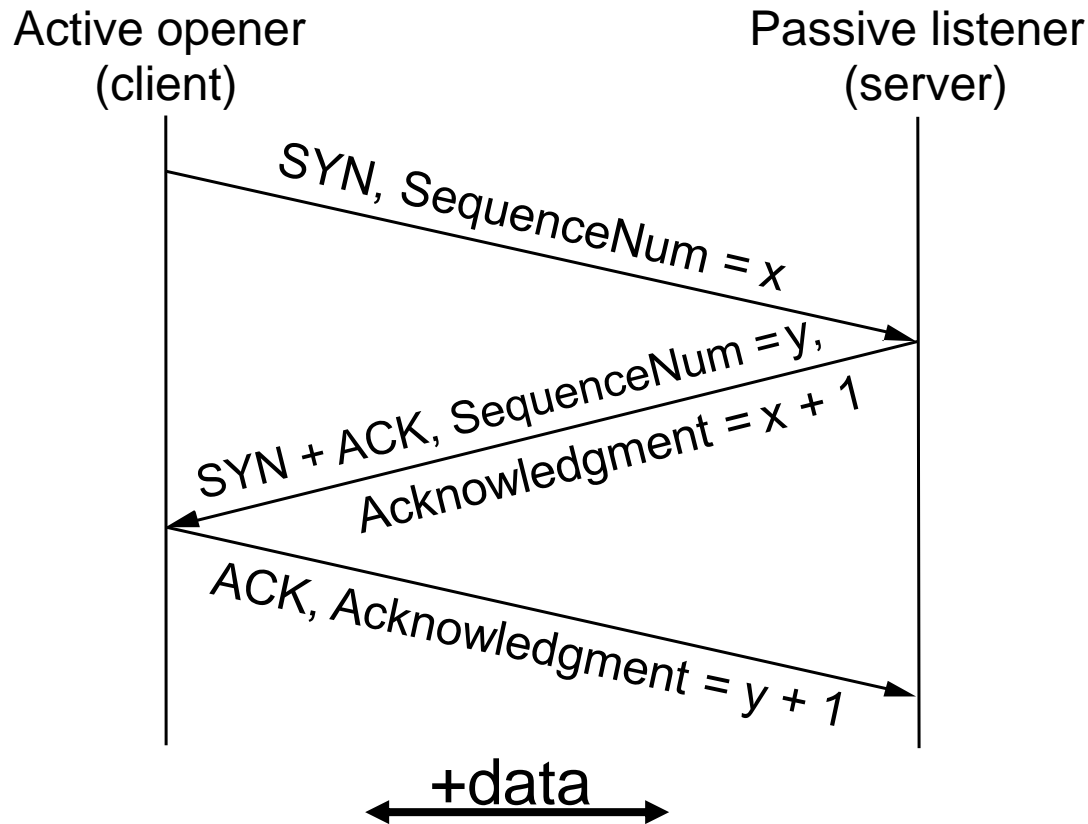
Source port		Destination port	
Sequence number			
Acknowledgement			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			

## TCP Header Fields

- Each connection identified with 4-tuple:
  - (SrcPort, SrcIPAddr, DsrPort, DstIPAddr)
- Sliding window + flow control
  - acknowledgment, SequenceNum, AdvertisedWinow
- Flags
  - SYN, FIN, RESET, PUSH, URG, ACK
- Checksum
  - pseudo header + TCP header + data

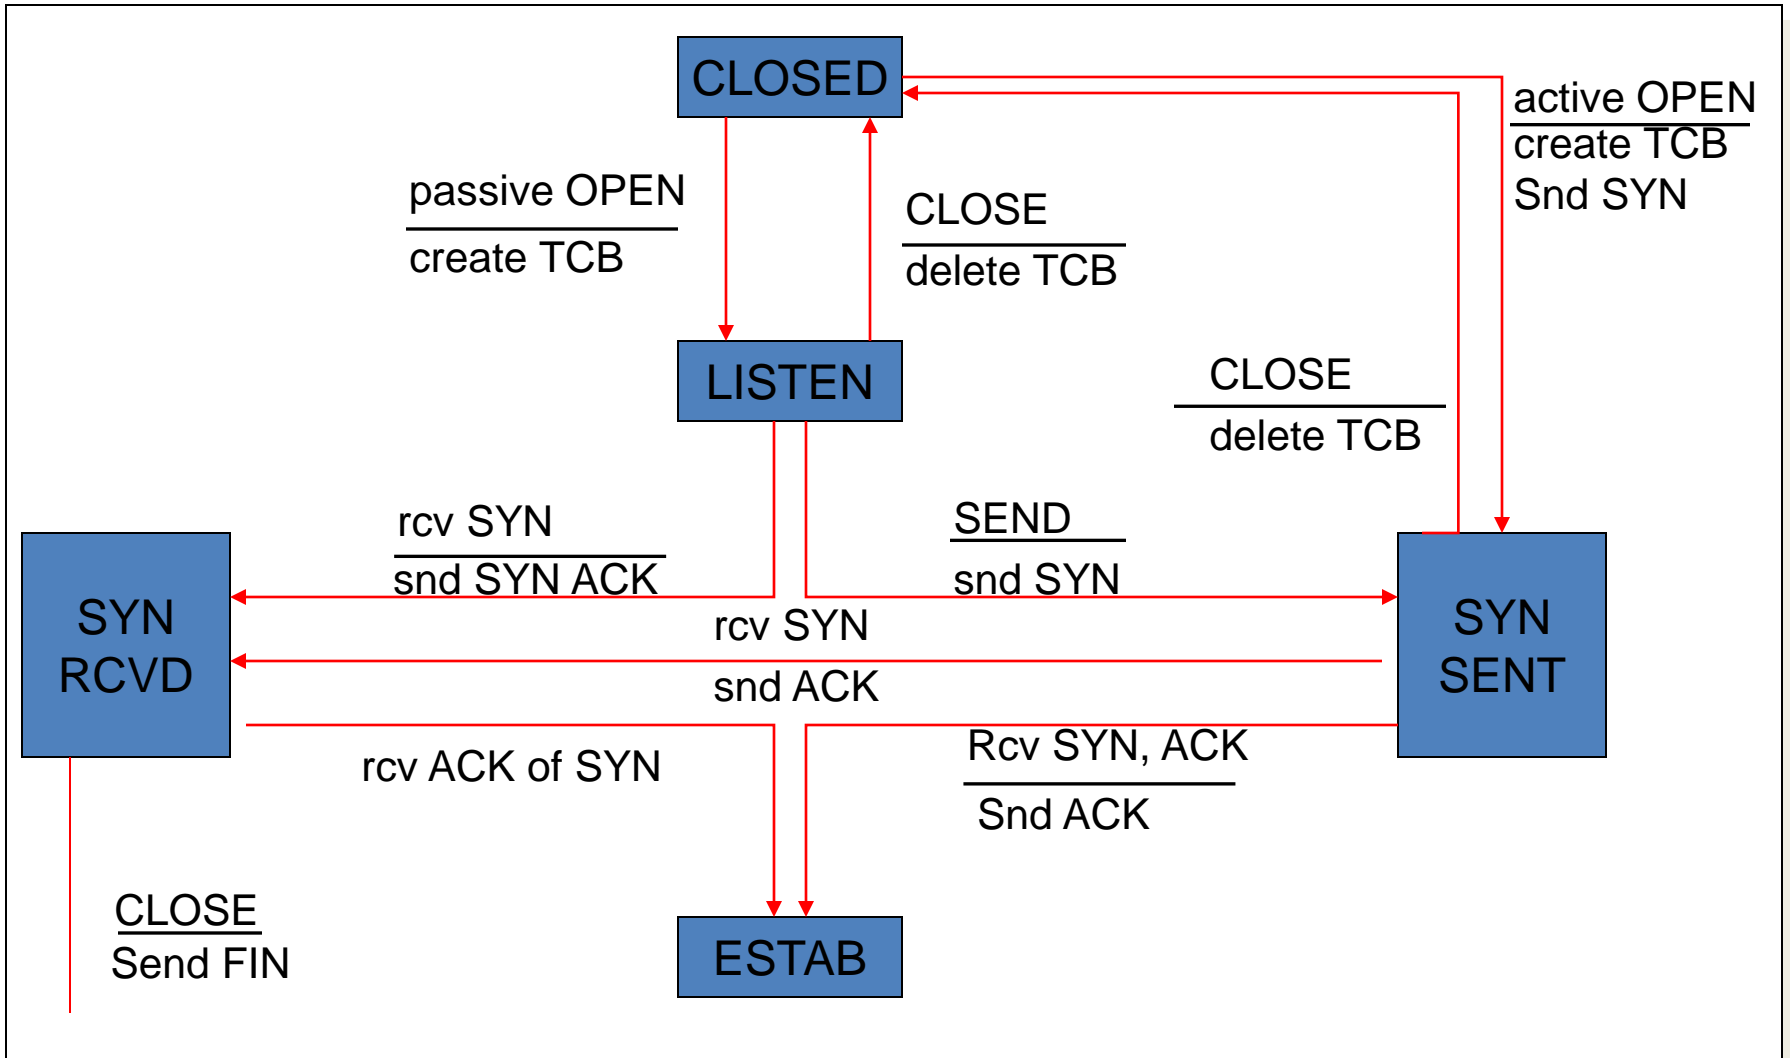
# Three-Way Handshake

- Opens both directions for transfer

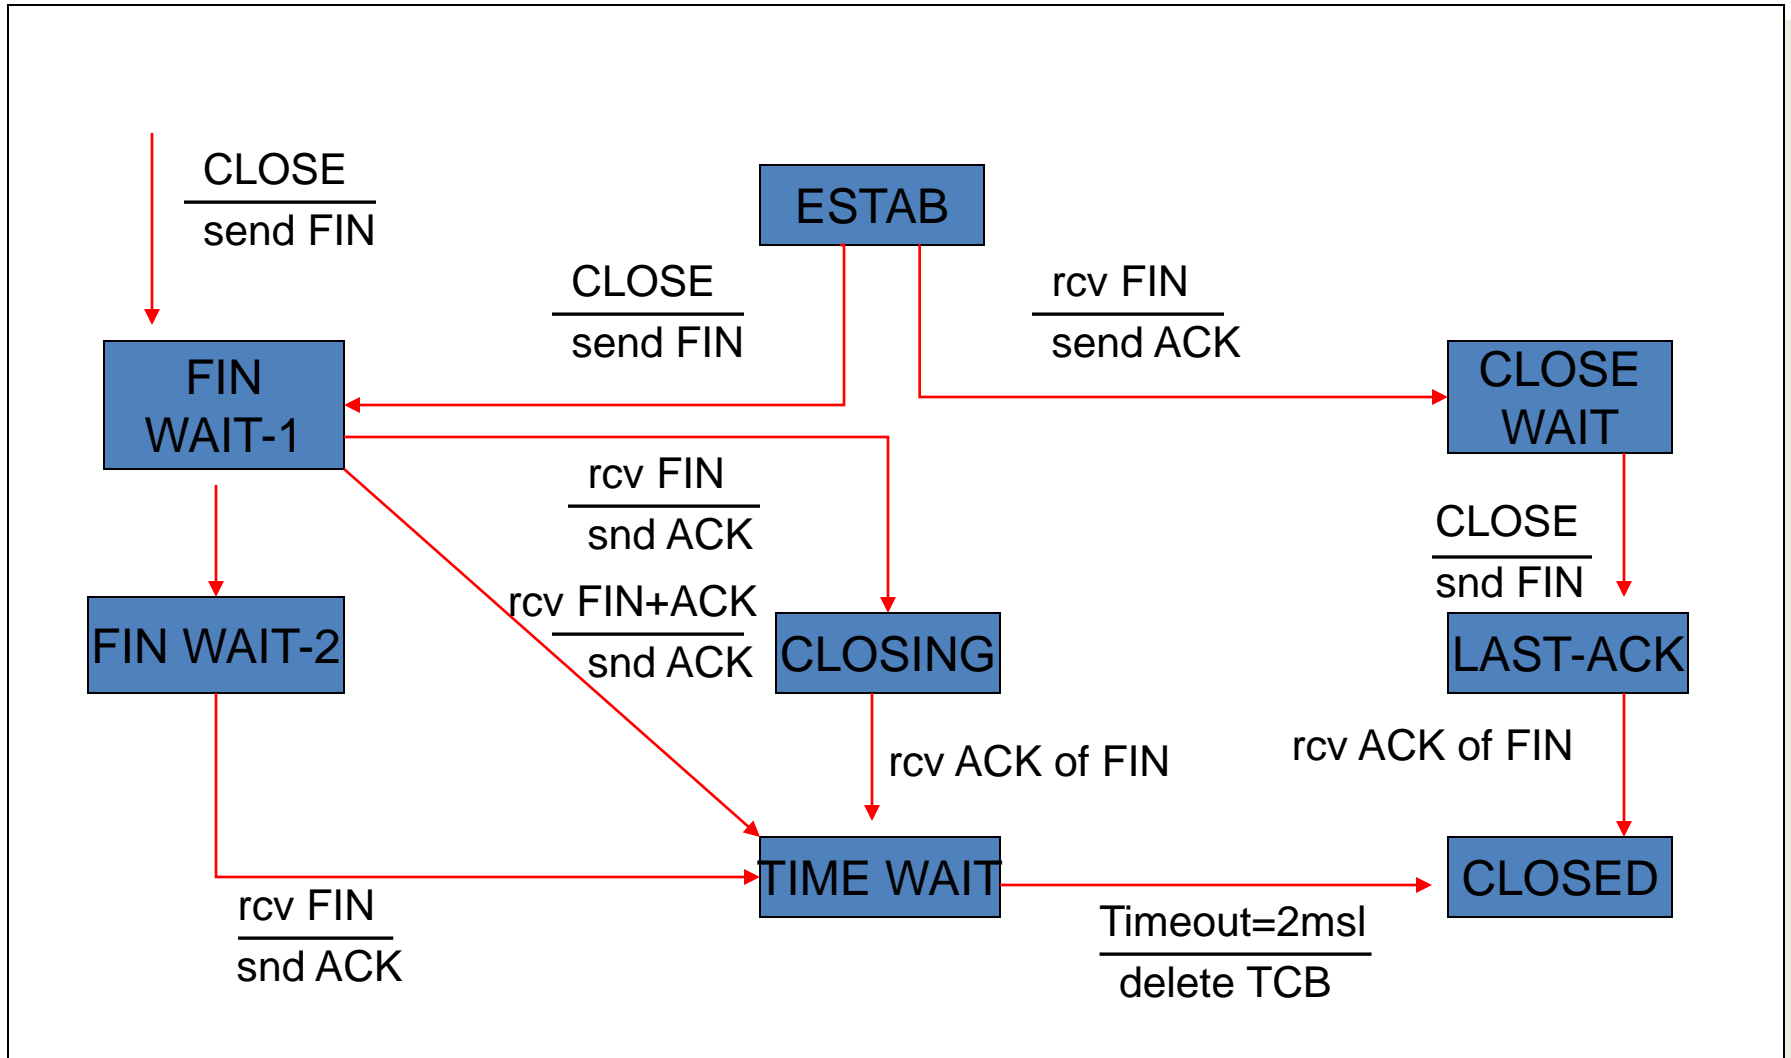




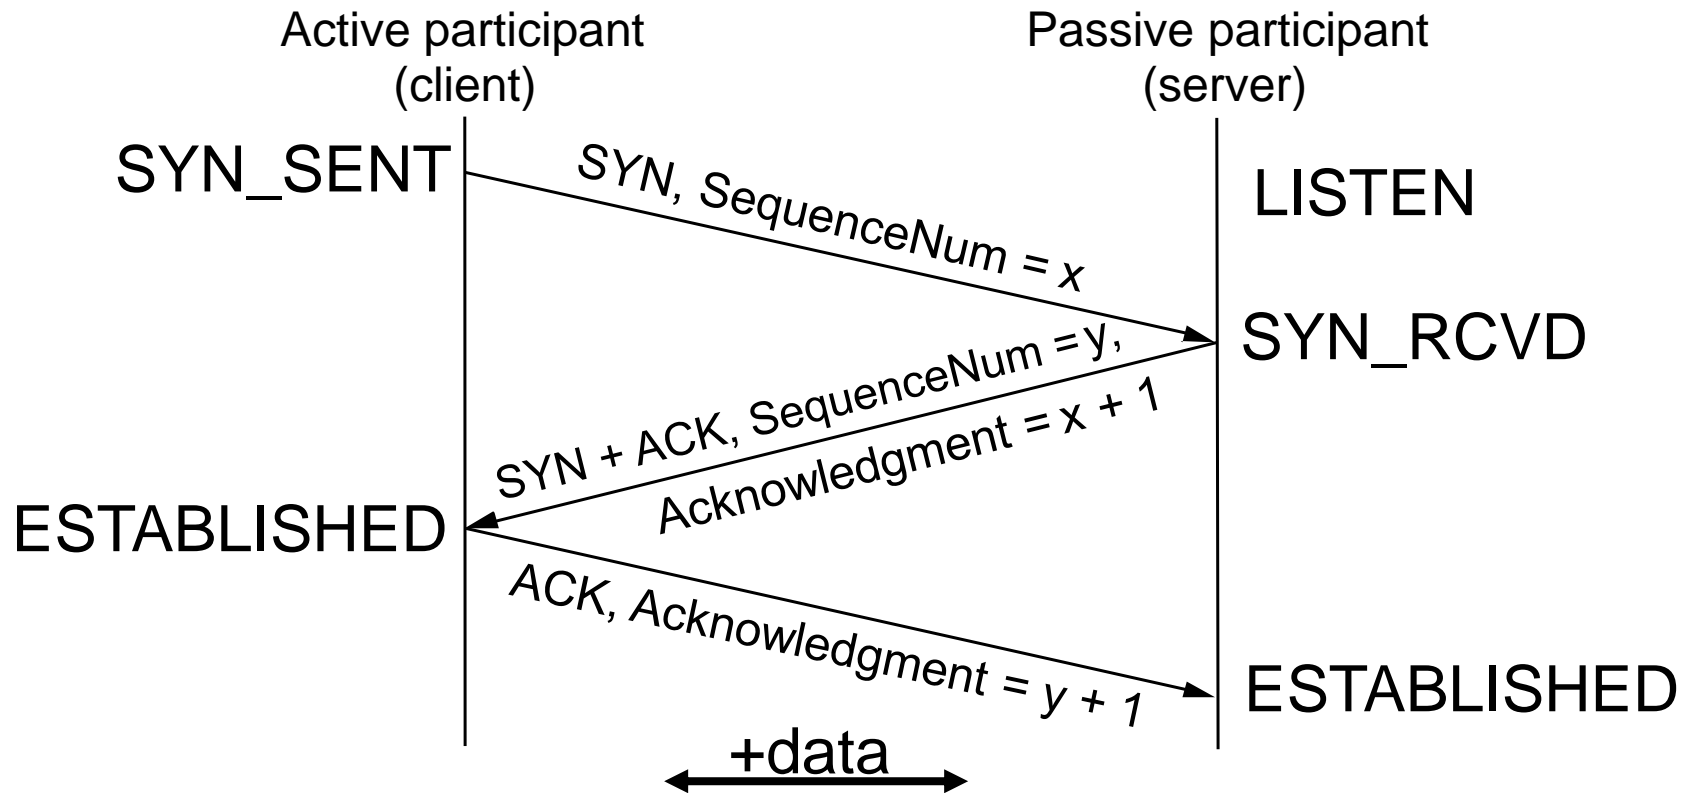
# TCP Connection Setup: FSM



# TCP Connection Tear-down: FSM



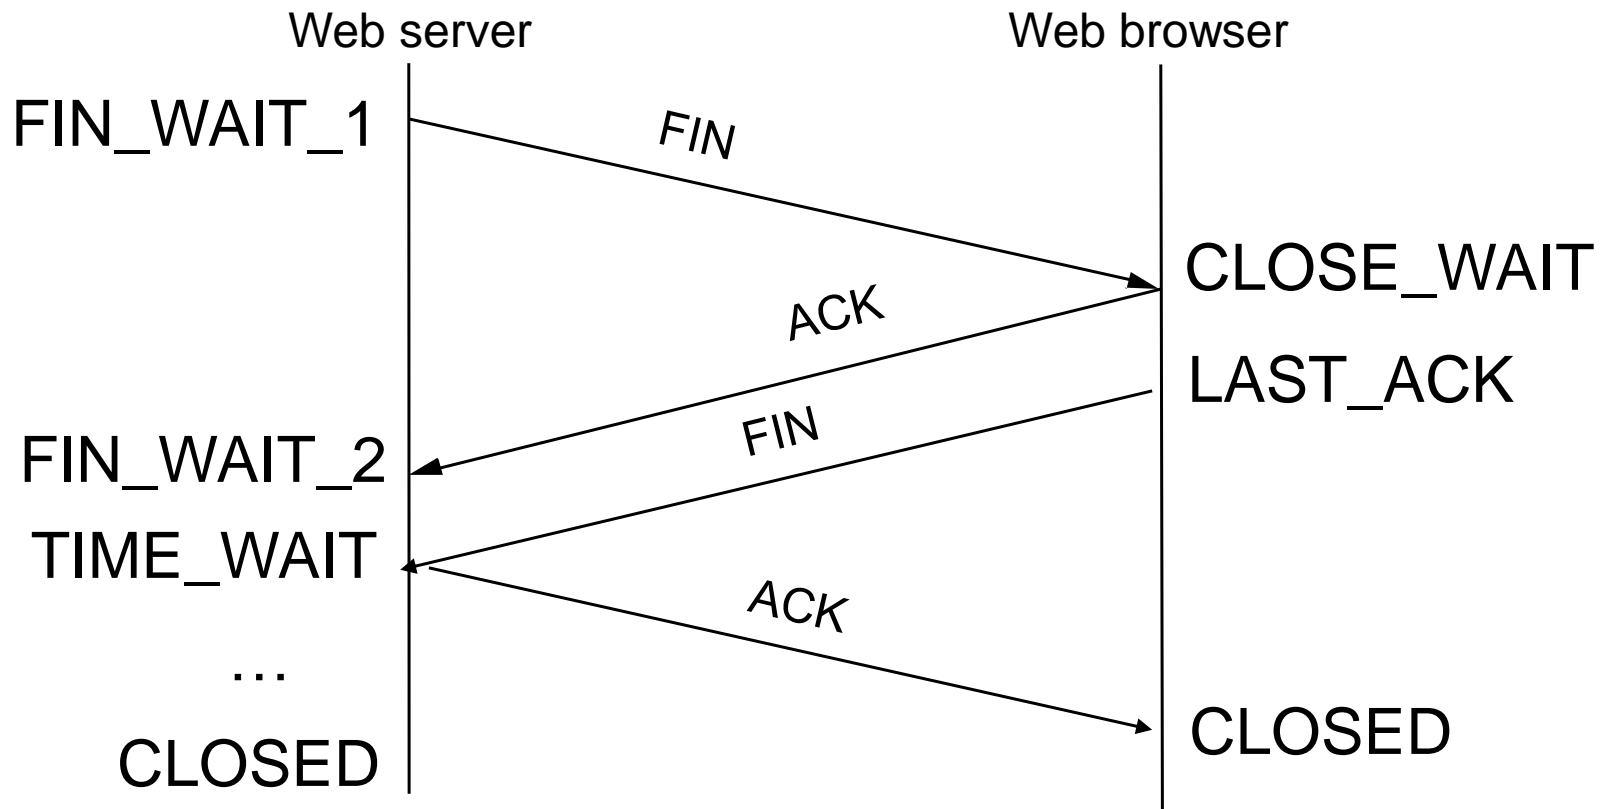
# Again, with States



# Connection Teardown

- Orderly release by sender and receiver when done
  - Delivers all pending data and “hangs up”
- Cleans up state in sender and receiver
- TCP provides a “symmetric” close
  - both sides shutdown independently

# TCP Connection Teardown




# The TIME\_WAIT State

- We wait 2MSL (two times the maximum segment lifetime of 60 seconds) before completing the close
- Why?
- ACK might have been lost and so FIN will be resent
- Could interfere with a subsequent connection

# Sequence Number Space

- Each byte in byte stream is numbered.
  - 32 bit value
  - Wraps around
  - Initial values selected at start up time
- TCP breaks up the byte stream in packets.
  - Packet size is limited to the **Maximum Segment Size**
- Each packet has a sequence number.
  - Indicates where it fits in the byte stream



# MSS

## ❑ Maximum Segment Size (MSS)

- ❑ Largest “chunk” sent between TCPs.
  - ❑ Default = 536 bytes. Not negotiated.
  - ❑ Announced in connection establishment.
  - ❑ Different MSS possible for forward/reverse paths.
  - ❑ Does not include TCP header
- What all does this effect?
  - Efficiency
  - Congestion control
  - Retransmission
- Path MTU discovery
  - Why should MTU match MSS?



# Silly Window Syndrome

- Problem: (Clark, 1982)
  - If receiver advertises small increases in the receive window then the sender may waste time sending lots of small packets
- Solution
  - Receiver must not advertise small window increases
  - Increase window by  $\min(\text{MSS}, \text{RecvBuffer}/2)$

# Nagel's Algorithm & Delayed Acks

- Small packet problem:
  - Don't want to send a 41 byte packet for each keystroke
  - How long to wait for more data?
- Solution: Nagel's algorithm
  - Allow only one outstanding small (not full sized) segment that has not yet been acknowledged
- Batching acknowledgements:
  - **Delay-ack timer:** piggyback ack on reverse traffic if available
  - 200 ms timer will trigger ack if no reverse traffic available

# Timeout and RTT Estimation

- Problem:
  - Unlike a physical link, the RTT of a logical link can vary, quite substantially
  - How long should timeout be ?
  - Too long => underutilization
  - Too short => wasteful retransmissions
- Solution: adaptive timeout: based on a good estimate of maximum current value of RTT

# How to estimate max RTT?

- $RTT = \text{prop} + \text{queuing delay}$ 
  - Queuing delay highly variable
  - So, different samples of RTTs will give different random values of queuing delay

# Round Trip Time and Timeout (II)

Q: how to estimate RTT?

- SampleRTT: measured time from segment transmission until ACK receipt
- SampleRTT will vary wildly
  - use several recent measurements, not just current SampleRTT to calculate “AverageRTT
  - **AverageRTT = (1-x)\*AverageRTT + x\*SampleRTT**
  - Exponential weighted moving average (EWMA)
  - Influence of given sample decreases exponentially fast; x = 0.1 or 0.2

## Setting the timeout

$$\text{Timeout} = \text{AverageRTT} + 4 * \text{Deviation}$$

$$\text{Deviation} = (1-x) * \text{Deviation} + x * |\text{SampleRTT} - \text{AverageRTT}|$$

# Karn's RTT Estimator

- Accounts for *retransmission ambiguity*
- If a segment has been retransmitted:
  - *Don't update RTT estimators during retransmission.*
  - **Timer backoff:** If timeout,  $RTO = 2 * RTO$  {exponential backoff}
    - Keep backed off time-out for next packet
  - Reuse RTT estimate only after one successful packet transmission